

Dgrammar and DPGrammar: Frontier Constrained Decoding for Diffusion Language Models

Jeng-Yue Liu¹, Wilson Zheng¹, Haoling Pu¹

¹Language Technologies Institute, Carnegie Mellon University
{buffettl,wilsonz,haolingp}@andrew.cmu.edu

Abstract

The leading constrained decoding method for discrete diffusion language models (dLLMs) reports about 99% validity on simple JSON benchmarks, but this does not transfer to harder schemas. On JSONSchemaBench medium dataset, stochastic lookahead-verify achieves only 78.9% validity, consuming 255.3 lookahead draws per instance and timing out on 61 instances. We address this in two steps. **Dgrammar** replaces stochastic lookahead with deterministic frontier masking, AIMD batching, and async mask overlap, raising validity to 84.9% and eliminating all timeouts. **DPGrammar** adds a Viterbi Dynamic Programming (DP) that jointly repairs violated token spans over grammar DFA¹ states, reaching **95.3%** validity with 2.11 mean resamples. Code: https://github.com/buffett0323/anlp_final.

1 Introduction

Many applications require model outputs to satisfy a formal grammar or schema: APIs that consume JSON, tools that compile code, pipelines that validate structured records. dLLMs such as LLaDA (Nie et al., 2025) and Dream (Ye et al., 2025) are attractive because they refine many tokens simultaneously under bidirectional context, but this parallel update breaks the monotonic prefix on which classical Finite-State Machine (FSM)² constrained decoding is built. Ensuring schema compliance without destroying the model’s joint proposals is therefore an open problem with direct deployment consequences.

The leading constrained decoding method for dLLMs is LAVE (Zhang et al., 2026). While LAVE achieves near-perfect validity on json-mode-eval

(Nous Research, 2024), our evaluation on the more complex JSONSchemaBench (Geng et al., 2025) medium reveals a performance gap, with validity decreasing to 78.9% on the $n=511$ instances for which LLGUIDANCE (Microsoft, 2025) accepts the schema due to the benchmark’s intricate nested structures and cross-field constraints. The gap is not incidental: LAVE verifies each frontier token by drawing K random completions and accepting only if at least one is grammar-valid. On simple schemas, random suffixes are likely to be valid; on schemas with joint multi-token constraints, the probability that a random suffix satisfies all constraints simultaneously drops sharply, causing false rejections, K -fold forward cost per token, and stochastic run-to-run instability.

Figure 1 concretizes the three regimes above: stochastic lookahead, greedy frontier remasking, and joint span repair on one running instance. We address the root cause in two steps. **Dgrammar** replaces stochastic sampling with deterministic frontier masking: grammar constraints use the incremental LLMatcher from LLGUIDANCE (Microsoft, 2025) at the first unresolved position only; committed spans advance the parser state without replaying from scratch; AIMD³ adaptive batching unmask a block per forward pass; violators are localized and resampled in-place; and computes mask asynchronously overlapping the GPU forward, where no random completions drawn. This engineering redesign raises validity to 84.9% and cuts mean wall time $3\times$. When a batch unmask produces tokens with inter-dependent joint violations (e.g., an enum mismatch at c and a regex-pattern violation at $s-1$ must be satisfied simultaneously), greedy per-position repair still falls short. **DPGrammar** adds a Viterbi dynamic program that jointly reas-

¹A Deterministic Finite Automaton (DFA) is a theoretical model of computation consisting of a fixed set of states and transitions that accepts or rejects a string of symbols by following a unique, predetermined path for each input.

²FSM refers to the automaton that tracks which tokens are valid at each position given the grammar/schema.

³Adapted from the Additive Increase/Multiplicative Decrease (AIMD) principle in TCP congestion control: B increases by 1 upon success and resets to 1 upon a grammar violation.




Figure 1: Motivating example from JSONSchemaBench medium instance o9875 (API route config). The schema requires $\text{method} \in [\text{"get"}, \text{"post"}, \text{"put"}, \dots]$ (lowercase enum) and each $\text{experienceGroupIds}$ item to match $^{\wedge}[\text{A-Fa-f0-9}]\{24\}\$$ (24-character hex, e.g. MongoDB ObjectId). (1) The diffusion model’s batch unmask produces "GET" (fails lowercase enum, position c) and "grp_abc" (fails 24-char hex pattern, position $s-1$); the violated span covers both. (2a) LAVE samples K random completions from c : satisfying both the enum *and* the exact-24 hex pattern simultaneously requires $P \approx (16/|V|)^{24} \approx 0$, no valid completion is ever drawn; LAVE times out at 120 s (255+ mean draws overall). (2b) Dgrammar forces "get" at c via LLMatcher masking (enum fixed), but the hex-pattern violation at $s-1$ requires 100 resamples to generate a valid 24-character string (84.9% overall validity). (2c) DPGrammar runs Viterbi over $[c, s)$ jointly, assigning both "get" and a valid 24-char hex ID in one pass, only 7 resamples on this instance (95.3% overall validity).

signs the entire violated span $[c, s)$ over grammar DFA states and top- K token candidates in one pass, pushing validity to 95.3% with mean resamples of 1.06.

Our contributions are: (i) identifying and quantifying LAVE’s benchmark gap, tracing it to stochastic sampling under complex joint constraints; (ii) a deterministic engineering redesign achieving 84.9% validity; (iii) a principled Viterbi repair layer achieving 95.3%, matching LAVE’s claimed number, now on a harder benchmark; and (iv) fine-grained timing instrumentation separating grammar-check from forward-pass cost.

2 Related Work

Discrete diffusion and parallel decoding. DLLMs such as LLaDA (Nie et al., 2025) and Dream (Ye et al., 2025) enable non-autoregressive (AR), any-order text generation through iterative masked refinement. This flexibility introduces the

incomplete prefix problem: unlike AR models, where a stable left-to-right prefix always exists for grammar-state tracking, dLLMs must predict tokens surrounded by still-masked or partially denoised positions, making it unclear where the parser frontier lies. Blockwise accelerators (e.g., DFlash (Chen et al., 2026), Spiffy (Agrawal et al., 2026)) reduce sampling steps but remain unconstrained, offering no structural or syntactic guarantees. Our work is orthogonal: we fix the backbone and study how to impose grammar constraints despite the incomplete prefix.

Constrained decoding in AR models. Tooling such as Outlines (Willard and Louf, 2023) and XGrammar (Dong et al., 2024) exploits a monotonic prefix to maintain FSM or parser states across generation steps. These interfaces do not transfer to diffusion because there is no single causal frontier until we define one.

Constrained decoding for diffusion. Mündler et al. (2025) propose the first CFG-constrained dLLM decoder; DINGO (Suresh et al., 2025) uses dynamic programming for syntactic correctness at higher per-step cost. LAVE (Zhang et al., 2026) introduces a lookahead-then-verify mechanism, sampling K random completions to decide whether to accept each frontier token, and has become the dominant baseline, reporting 99% syntactic validity on json-mode-eval (Nous Research, 2024). We show this number is benchmark-specific: LAVE’s stochastic sampling fails to cover joint constraint spaces on harder schemas, dropping to 78.9% on JSONSchemaBench (Geng et al., 2025) medium on the LLGUIDANCE-compatible instances, and its reliance on random completions risks false rejection of genuinely valid tokens. Plan-style decoders (e.g., PVF (Li et al., 2026)) emphasize structure-aware planning; we stay closer to token-level masking with an explicit incremental matcher.

Sparse transition structures and DP cost. STATIC (Su et al., 2026) shows that sparse DFA transition structures can be flattened into Compressed Sparse Row (CSR) matrices, enabling $O(K)$ coalesced lookups per step, where K is the grammar branch factor, rather than scanning the full vocabulary $|V|$. We adapt this principle to DP-Grammar: instead of scoring all $|V|$ tokens at each DP position, we restrict the Viterbi search to top- K candidates per position, reducing per-step transition cost from $O(|V|)$ to $O(K)$ while preserving coverage of the most probable grammar-feasible paths.

3 Data

We evaluate on jsb_medium, a JSON Schema benchmark split used in our dLLM evaluation harness, aligned with the JSON schema structured output setting discussed in broader benchmarks (Geng et al., 2025). Each instance provides a natural language style prompt, a schema string, and evaluation extracts structured output from the model generation region.

Scale and splits. We merge JSONL logs from nine Modal shards (seed 0, $T=128$, offsets 0, 66, ..., 528) on instance_id. The driver originally targets $n=586$ instances. Dgrammar and DP-Grammar build constraints with LLGUIDANCE (Microsoft, 2025); when JSON Schema compilation fails (unsupported keywords, draft incompatibili-

ties, circular references, or malformed specs), no matcher is created and that instance cannot be run. This removes 75 instances. For a controlled comparison, we evaluate all three methods on the same remaining $n=511$ instances, we exclude those 75 cases for LAVE as well, so headline metrics refer to an identical instance list even though LAVE’s own codebase can ingest a wider class of schemas.

Preprocessing. Instances without a usable schema string are skipped in the driver. No additional manual annotation is performed; validity is judged by the benchmark’s extraction and schema check, surfaced as the boolean valid in result JSON lines.

4 Method

4.1 Baselines reproduced

Our primary external baseline is LAVE (Zhang et al., 2026), reimplemented to operate with our LLaDA-8B-Instruct backbone and jsb_medium loader under the same diffusion step budget $T=128$ where applicable. This isolates the constraint algorithm rather than changing the base model. We additionally compare Dgrammar against DP-Grammar on identical instances.

4.2 Dgrammar: proposed deterministic frontier stack

From the engineering perspective, we target two coupled issues in constrained decoding for masked diffusion LLMs: (i) hard masking or lookahead over whole spans can distort the model’s joint proposals across positions; (ii) per-step grammar checks and mask construction can add latency on top of already expensive diffusion forwards. We avoid stochastic suffix sampling entirely and instead enforce the schema only at the decoding frontier, while reusing parser state and hiding much of the constraint work behind the next GPU forward.

Incremental grammar backbone. We build on LLGUIDANCE (Microsoft, 2025), which exposes an incremental finite-state view of the schema: each newly committed token advances a persistent automaton state, so we never re-scan the full prefix from the prompt after every block update.

Frontier-only constraints. Constraint enforcement applies at the first position that is still undecided. We derive the allowed next-token set there from the current automaton state; elsewhere, positions may remain masked while we read committed

text left to right. Logits stay unconstrained off the frontier, and we avoid locking down an entire parallel row before the batch of unmask decisions is mutually consistent.

State across denoising steps. As spans stabilize across diffusion iterations, we carry the automaton forward rather than re-initializing from the prompt on every check, which keeps repeated validity tests cheap.

Adaptive commit (AIMD). When consumption succeeds we grow how many tokens we commit per step, up to $k_{\max}=8$; when a conflict appears we shrink the batch, locate the first offending position, repair it, and only then widen the commit window again.

Violator remask and in-place resampling. We mask only the offending token, set its logit to $-\infty$ in the current forward output, and walk successive argmax alternatives on that same tensor so many repairs do not require another GPU forward.

Early termination and asynchronous overlap. We halt once the automaton indicates a complete, schema-valid string, skipping superfluous late diffusion steps. Separately, we prepare the next frontier’s allowed-token mask on a CPU worker while the GPU executes the following forward, so mask construction overlaps forward time and reduces wall-clock constraint overhead.

4.3 DPGrammar: segment DP on top of Dgrammar

DPGrammar reuses the same denoising and frontier schedule as Dgrammar; the only difference is how we repair a rejected continuation. Greedy single-position remasking can leave correlated errors when a batch unmask reveals several dependent tokens at once. DPGrammar therefore treats the violated interval $[c, s)$, where c is the first offending position and $s \triangleq \text{find_constraint_end}(c)$ is the nearest bracket-balanced, grammar-consistent boundary, as a single joint decision. We run a Viterbi DP over that segment, combining grammar automaton states with the $K=100$ highest-scoring vocabulary entries at each position under the current forward. The DP searches for a replacement sequence that restores a consistent parse; if no feasible path exists within that top- K beam, we fall back to Dgrammar’s greedy remask. Figure 2 sketches the shared stack and where this segment-level re-

pair attaches. The full pseudocode of the algorithm is given in Appendix A.

4.4 Completion after generation

Both systems share driver side repair after the diffusion phase: grammar guided greedy autocomplete, masked extension, grammar only closing, and minimal JSON synthesis from the schema when required.

4.5 Implementation and compute

Code builds on LLGUIDANCE (Microsoft, 2025), and the constrained diffusion CD4dLLM style bindings (Mündler et al., 2025; Zhang et al., 2026); benchmarks run on Modal with $9 \times$ NVIDIA A100 GPUs in our experiments. Backbone is GSAI ML LLaDA-8B-Instruct⁴. Block length 32 for timed async runs, remasking low confidence, temperature 0.2 in the driver. Full hyperparameters are listed in Appendix D.

4.6 Experimental design and metrics

Table 1: Experimental Metrics and Definitions

Metric	Description
Schema Validity	Primary outcome; binary (valid) per instance.
Wall Time	End-to-end driver time (mean seconds per instance).
Constraint OH %	Driver-reported grammar and mask work as a fraction of total time.
Eff. Constraint %	Total constraint cost including mask-wait and autocomplete (Dgrammar/DPGrammar).
Resamples	Count of remasking or lookahead verification draws (LAVE).
Per-token Time	Average constraint and total work normalized over decoded length.

5 Results

Table 2 reports aggregate metrics on the $n=511$ instances retained after excluding schemas that LLGUIDANCE cannot compile. We discuss results in the two-step framing of our approach.

LAVE vs. Dgrammar (engineering step). On this matched set, LAVE reaches 78.9% validity with 255.32 mean lookahead draws per instance and 39.48 s mean wall time. Dgrammar replaces stochastic sampling with deterministic frontier masking: validity rises to 84.9%, mean wall time falls to 13.14 s ($\approx 3 \times$ faster than LAVE), and mean resamples fall to 32.43. The engineering redesign,

⁴<https://huggingface.co/GSAI-ML/LLaDA-8B-Instruct>




Figure 2: **Left:** one denoising step under Dgrammar (blue) and the DPGrammar branch (orange). **Right:** why joint repair matters. **(a)** Greedy fixes one position at a time: after forcing the correct token at c , the constraint at $s-1$ may still be violated, triggering another resample, which can itself cascade, 32.43 mean resamples per instance. **(b)** The Viterbi DP sweeps the span $[c, s)$ once, tracking grammar DFA states and selecting the highest-log-prob jointly grammar-valid token sequence, 1.06 mean resamples per instance.

incremental parser state, AIMD batching, async overlap, accounts for this gain without a new theoretical primitive.

Dgrammar vs. DPGrammar (principled step). Dgrammar’s residual failures stem from greedy per-position repair that misses joint constraint structure (e.g., fixing a key but leaving an incompatible value type in the same span). DPGrammar adds `dp_fix_prefix`: a Viterbi search over the full violated span under grammar DFA states. Validity rises to 95.3%, mean resamples collapse from 32.43 to 1.06, while mean wall time moves from 13.14 s to 14.96 s, a modest increase relative to the validity gain. As shown in Table 2, effective constraint time as a fraction of wall clock drops sharply. The principled extension closes what the engineering step cannot: multi-token joint violations that greedy remasking iterates through slowly.


Qualitative observations. Error modes for the remaining 4.7% DPGrammar failures include completion-phase limits and schemas where minimal-JSON fallback is brittle. Detailed case studies are left to appendix space.

6 Discussion

The LAVE–Dgrammar gap shows that deterministic frontier parsing is a strong engineering solution: it eliminates sampling variance and timeout risk with no new theoretical machinery. The Dgrammar–DPGrammar gap, in turn, shows that greedy token-by-token repair has a structural limit, joint violations spanning multiple positions call for a principled combinatorial search, not more engineering tuning. The two steps are complementary: the engineering stack is prerequisite infrastructure; the DP layer addresses what greedy repair leaves unresolved. Empirically, DPGrammar attains the highest validity and the lowest resample churn, while Dgrammar remains an attractive option wherever the overhead of segment DP is unwelcome.

These results should be read with several caveats. Effective constraint % and mask timings are not equally interpretable across methods (LAVE lacks the same breakdown). DPGrammar’s mean wall time is slightly above Dgrammar’s despite far fewer resamples, reflecting DP work on violation spans. DPGrammar still falls short of 100% validity, so driver-side completion and fallback paths remain load-bearing rather than cosmetic.

LAVE: stochastic lookahead-then-verify



Dgrammar: deterministic frontier masking




Figure 3: One denoising step under schema {"count":int, "label":str}. **LAVE** (Zhang et al., 2026) (left) samples K random completions at each frontier token, accepting it only when ≥ 1 suffix is valid—stochastic and $K\times$ more expensive per token, hitting the 120 s per-instance timeout on hard schemas. **Dgrammar** (right) runs compute_mask on the incremental FSM state and unmask an entire block under that mask in one forward pass, overlapped with the next mask computation; the result is deterministic.

	LAVE	Dgrammar	DPGrammar
Skipped (grammar-invalid)	75	75	75
Evaluated (n)	511	511	511
Valid	397 / 511	434 / 511	487 / 511
Validity (%)	77.7%	84.9%	95.3%
Timeouts (>120 s)	68	0	0
Mean resamples [§]	233.87	32.43	2.11
Mean time (s)	41.12	13.38	16.14
Median time (s)	27.34	13.65	15.61
P95 time (s)	120.00	20.93	29.36
Max time (s)	120.05	33.45	88.17

Table 2: jsb_medium results (seed 0, $T=128$ steps). All three methods are evaluated on the same $n=511$ matched instances; 75 instances are skipped by all methods due to grammar-invalid schemas (LAVE: immediate parse error; Dgrammar/DPGrammar: silently skipped by LLGUIDANCE). [§] LAVE resamples count lookahead verification draws; Dgrammar/DPGrammar resamples count token re-mask events.

7 Conclusion and Future Work

We studied grammar constrained decoding for dLLMs in two steps. Dgrammar is an engineering redesign of the constraint loop that raises validity from 77.7% (LAVE) to 84.9% and cuts mean wall time from 41.12 s to 13.38 s without any new theoretical primitive. DPGrammar is the principled extension: Viterbi-based joint repair over violated spans pushes validity to 95.3% and collapses resamples from 32.43 to 2.11, addressing multi-token joint constraints that greedy repair cannot resolve efficiently. Broader takeaway: the engineering step provides the infrastructure; the principled step provides the correctness guarantee; both are needed for near-perfect structured generation under diffusion.

Future work includes extending to additional JSON splits or code grammars; and exploring

tighter integration with sparse automata acceleration (Su et al., 2026) for very large vocabularies.

8 Limitations

Dataset scope. All results are restricted to the $n=511$ LLGUIDANCE-compatible instances of jsb_medium (see Appendix E). We evaluate on a single benchmark split and a single backbone (LLaDA-8B-Instruct); generalization to other JSON splits, code grammars, or larger models remains untested.

Validity vs. content quality. Grammar-constrained decoding guarantees *structural* validity but cannot improve semantic content quality. We observe that LLaDA-8B sometimes produces degenerate yet valid completions, empty arrays (`[]`), empty objects (`{}`), or single-character

strings ("x"), particularly for schemas with additionalProperties: false or deeply nested array items. These outputs pass the JSON Schema validator and are counted as valid in Table 2, but they carry little semantic content. This is a property of the underlying diffusion model rather than the constraint method: LAVE and Dgrammar exhibit the same pattern on the same instances.

DP approximation. DPGrammar explores only the top- K tokens per position ($K=100$), so a globally optimal grammar-valid token outside the top- K can be missed. The constraint boundary heuristic (find_constraint_end) uses bracket-depth tracking on original tokens; it may over- or under-extend the repair span on unusual schema structures. When the DP finds no valid path it falls back to remasking the violator, inheriting LAVE-style retry cost for those instances.

Hyperparameter sensitivity. We fix $T=128$, block length 32, and temperature 0.2 across all methods. Different settings could shift absolute numbers; relative rankings may be more stable. Top- K for the DP and the timeout threshold were not tuned beyond initial exploration.

9 Ethical Considerations

Structured generation can reduce malformed outputs in downstream automation, but valid JSON is not always semantically safe or fair: schemas can encode biased or incomplete coverage of user intent. Models and grammars should be reviewed for misuse in high stakes settings (medical, financial). Our benchmark instances originate from public style schema tasks; they may not represent all languages or cultures. Researchers and deployers should combine syntactic guarantees with content moderation and policy controls. We use open backbone weights and public style evaluation data; users of our code should respect model licenses and avoid deploying unreviewed systems in sensitive domains without human oversight.

References

- Sudhanshu Agrawal, Rishiek Garrepalli, Raghav Goel, Mingu Lee, Christopher Lott, and Fatih Porikli. 2026. [Spiffy: Multiplying diffusion llm acceleration via lossless speculative decoding](#). *Preprint*, arXiv:2509.18085.
- Jian Chen, Yesheng Liang, and Zhijian Liu. 2026. [Dflash: Block diffusion for flash speculative decoding](#). *Preprint*, arXiv:2602.06036.
- Yixin Dong, Charlie F Ruan, Yaxing Cai, Ruihang Lai, Ziyi Xu, Yilong Zhao, and Tianqi Chen. 2024. Xgrammar: Flexible and efficient structured generation engine for large language models. *Proceedings of Machine Learning and Systems 7*.
- Saibo Geng, Hudson Cooper, Michał Moskal, Samuel Jenkins, Julian Berman, Nathan Ranchin, Robert West, Eric Horvitz, and Harsha Nori. 2025. [Json-schemabench: A rigorous benchmark of structured outputs for language models](#). *Preprint*, arXiv:2501.10868.
- Miao Li, Hanyang Jiang, Sikai Cheng, Hengyu Fu, Yuhang Cai, Baihe Huang, Tinghan Ye, Xuanzhou Chen, and Pascal Van Hentenryck. 2026. [Plan, verify and fill: A structured parallel decoding approach for diffusion language models](#). *Preprint*, arXiv:2601.12247.
- Microsoft. 2025. [LLGuidance: Fast structured outputs for LLMs](#). GitHub repository.
- Niels Mündler, Jasper Dekoninck, and Martin Vechev. 2025. [Constrained decoding of diffusion llms with context-free grammars](#). *Preprint*, arXiv:2508.10111.
- Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-Rong Wen, and Chongxuan Li. 2025. [Large language diffusion models](#). *Preprint*, arXiv:2502.09992.
- Nous Research. 2024. [json-mode-eval](#). Hugging Face Datasets.
- Zhengyang Su, Isay Katsman, Yueqi Wang, Ruining He, Lukasz Heldt, Raghunandan Keshavan, Shao-Chuan Wang, Xinyang Yi, Mingyan Gao, Onkar Dalal, Lichan Hong, Ed Chi, and Ningren Han. 2026. [Vectorizing the trie: Efficient constrained decoding for llm-based generative retrieval on accelerators](#). *Preprint*, arXiv:2602.22647.
- Tarun Suresh, Debangshu Banerjee, Shubham Ugare, Sasa Misailovic, and Gagandeep Singh. 2025. [Dingo: Constrained inference for diffusion llms](#). *Preprint*, arXiv:2505.23061.
- Brandon T Willard and Rémi Louf. 2023. Efficient guided generation for large language models. *arXiv preprint arXiv:2307.09702*.
- Jiacheng Ye, Zhihui Xie, Lin Zheng, Jiahui Gao, Zirui Wu, Xin Jiang, Zhenguo Li, and Lingpeng Kong. 2025. [Dream 7b: Diffusion large language models](#). *Preprint*, arXiv:2508.15487.
- Yitong Zhang, Yongmin Li, Yuetong Liu, Jia Li, Xiaoran Jia, Zherui Li, and Ge Li. 2026. [Lookahead-then-verify: Reliable constrained decoding for diffusion llms under context-free grammars](#). *Preprint*, arXiv:2602.00612.

A DPGrammar Algorithm (dp_fix_prefix)

Algorithm 1 gives the full pseudocode for the Viterbi segment repair used by DPGrammar. Given a violated span $[c, s)$ detected during generation, the algorithm maintains a dictionary of live DFA states, each mapped to its highest-scoring path (matcher clone and cumulative log-probability). At each position it probes the top- K vocabulary entries via rollback, no cloning, and records the best (prev_state, token) pair for each reachable DFA state (Viterbi merge). After the loop, it clones one matcher per surviving state (Phase 2) and records the backtrack pointers. Backtracking from the best final state yields the minimal set of token replacements needed to make the span grammar-valid.

Two design choices keep the algorithm efficient. First, DFA states are identified by `bytes(matcher.compute_logit_bias())`, a compact proxy for the automaton node; this bounds the number of active states by the grammar’s DFA size ($\sim 20\text{--}200$ for typical JSON schemas) rather than by K^{depth} . Second, matcher clones are created only once per surviving state per step (Phase 2), not once per (state, candidate) pair; cloning cost is therefore $O(|\text{winners}|) \leq O(\text{DFA size})$ per position rather than $O(K \times |\text{states}|)$.

The span $[c, s)$ is bounded by `find_constraint_end`, which probes forward from c using the grammar automaton and tracks bracket depth in the original token sequence to avoid stopping inside an unclosed array or object. The DP is thus applied to a narrow window around the violated constraint rather than the full remaining sequence.

B Case Study: Instance o12610

We examine instance o12610 from `jsb_medium`, a representative case where LAVE completely fails while DPGrammar produces a clean, fully valid output in a single pass with zero resamples.

Schema. The schema describes a physical Wi-Fi radio device configuration (OpenWRT/LuCI format). The critical constraints are enum fields on `htmode` and `hwmode`:

```
"htmode": { "type": "string",
  "enum": [null, "HT20", "HT40+", "HT40-",
    "VHT20", "VHT40", "VHT80"] },
"hwmode": { "type": "string",
  "enum": ["11b", "11g", "11a"] }
```

Additional fields include `channel` (integer, range 1–165), `id` (string), `short_gi_20/40/80` (boolean), and `rx_stbc/tx_stbc` (integer, 0–1).

Why LAVE fails (timeout, 120 s, no output).

LAVE’s stochastic lookahead commits a token only if all K randomly sampled completions yield a grammar-valid continuation. For enum fields, this requires a random completion to land on exactly one of a small set of technical strings (e.g. "VHT20", "11a") drawn from a natural-language vocabulary. The probability is near zero: LLaDA-8B assigns negligible probability mass to "VHT20" or "HT40-" as free continuations, so every lookahead batch fails the grammar check. LAVE cannot commit even the first enum-constrained token and times out after 120 s without producing any output.

DPGrammar (valid, 0 resamples, 6.7 s).

Frontier masking restricts the vocabulary at each position to the set of tokens permitted by the grammar automaton. For the `htmode` field, only the four tokens spelling out a valid enum value are unmasked; the model samples from this constrained set directly. Because the automaton guides every token choice, no violation occurs and the Viterbi DP is never triggered (0 resample events, 0 DP calls). The output is complete and semantically coherent:

```
{
  "channel": 60,
  "htmode": "VHT20",
  "hwmode": "11a",
  "id": "pci-wifi-1",
  "rx_stbc": 1,
  "short_gi_20": true,
  "short_gi_40": true,
  "short_gi_80": false,
  "tx_stbc": 1
}
```

All enum constraints are satisfied (`htmode` $\in$ allowed set, `hwmode` = "11a"), the channel is within range, and boolean/integer fields are correctly typed. Total wall time is 6.7 s versus LAVE’s 120 s timeout, an $18\times$ speedup on this instance.

Summary. This instance illustrates the general failure mode of stochastic lookahead on enum-constrained fields: the probability of a random natural-language completion matching a specific technical string (e.g. "VHT20") is effectively zero, so LAVE never commits. Frontier masking (shared by Dgrammar and DPGrammar) sidesteps this entirely by restricting the model’s sampling distribution to the valid set at each position. On instances where frontier masking prevents all viola-

tions, DPGrammar’s Viterbi repair is never invoked, the method degrades gracefully to pure frontier-masked generation at the speed of a single unconstrained diffusion pass.

C Case Study: Format-Constrained String Fields (format: "uuid")

A recurring qualitative difference between Dgrammar and DPGrammar arises on schemas whose string fields carry *format* constraints (e.g., "format": "uuid", "format": "uri"). These constraints are enforced at the grammar level by LLGUIDANCE, which compiles them into the token automaton. The three methods therefore see different levels of constraint enforcement and produce noticeably different outputs.

Illustrative schema. Consider a `links` array whose items have a `type` string and a `target` string constrained to UUID format:

```
"links": {
  "type": "array",
  "items": {
    "properties": {
      "type": { "type": "string" },
      "target": { "type": "string",
                  "format": "uuid" }
    }
  }
}
```

The compiled grammar automaton restricts `target` to the pattern `[0-9a-f]{8}-[0-9a-f]{4}-... (RFC 4122 UUID)`, blocking all other strings.

Observed outputs across methods.

Dgrammar. Frontier masking restricts each token to the automaton’s valid set, but when the model generates a high-confidence non-UUID string (e.g. "12000e1TestCaseCanceledEvent"), each individual token may locally satisfy the grammar’s prefix condition while the *complete* string does not match the UUID pattern. Because Dgrammar’s validity check only tests whether an EOS token is present with no mask before it, it does not re-run a JSON Schema validator, the output is reported as `valid=True` even though `target` violates `format: "uuid"`. The output appears semantically natural but is in fact schema-invalid.

DPGrammar. The Viterbi DP operates on spans $[c, c+s)$ and selects the token sequence of maximal log-probability that is accepted by the grammar automaton. For a UUID field, the automaton enforces the full hexadecimal-and-hyphen pattern, so

the DP produces a well-formed UUID. The value "123e4567-90ab-12cd-abef-123456789000" is the model’s highest-probability completion of the UUID pattern given the prompt context, a *template UUID* that is grammatically correct but semantically generic. DPGrammar’s output is therefore schema-valid (unlike Dgrammar’s) but less contextually specific than one might hope.

LAVE. LAVE uses stochastic lookahead to verify token proposals. For UUID fields the lookahead must find a random completion that forms a valid UUID; if the model assigns low probability to hex digits, LAVE may fail repeatedly and fall back to a retry or timeout. On instances with many format-constrained fields, LAVE’s acceptance rate drops and wall time increases significantly.

Key takeaway: EOS validity $\neq$ schema validity. This case study illustrates why the valid flag (EOS/mask-based) is an *insufficient* proxy for correctness. Dgrammar reports `valid=True` on outputs that fail `format: "uuid"` validation, while DPGrammar’s schema-valid UUID is assigned the same `valid=True` flag. A schema validator or the *syntactic@k* metric (Section E) is required to distinguish the two cases. Conversely, DPGrammar’s correctly-formatted but semantically generic UUID highlights that grammar enforcement is a necessary but not sufficient condition for *functional* correctness: the generated UUID is unlikely to match any reference value, so *functional@k* (exact match) will be zero regardless of syntactic validity.

D Hyperparameters

All experiments use LLaDA-8B-Instruct as the backbone, run on Modal with NVIDIA A100 GPUs. Table 4 lists the shared settings across all three methods.

E Dataset and Instance Filtering

We evaluate on the `Github_medium` split of JSON-SchemaBench, which contains 586 instances. Of these, 75 are excluded from all comparisons because LLGUIDANCE rejects their schemas at compile time: LAVE immediately returns `valid=false` with the error "Prompt does not conform to grammar"; Dgrammar and DPGrammar silently skip them via the same LLGUIDANCE guard. Since all three methods behave identically on these 75 instances, including them would inflate

Method	type	target	Schema-valid?
Dgrammar	"self"	"12000e1TestCaseCanceledEvent"	No — violates format: "uuid"
DPGrammar	"relatedLink"	"123e4567-90ab-12cd-abef-123456789000"	Yes — well-formed UUID
LAVE	(varies / timeout)	(varies / timeout)	Depends on lookahead success

Table 3: Representative outputs on a jsb_medium instance with format: "uuid" on the target field.

Hyperparameter	Value
Diffusion steps T	128
Random seed	0
Block length	32
Remasking strategy	Low-confidence
Temperature	0.2
Top- K (DPGrammar DP)	100
Max resample attempts	500 (LAVE), 200 (ours)
Timeout threshold	120 s

Table 4: Shared hyperparameters for all benchmark runs on jsb_medium.

LAVE’s failure count without providing a meaningful comparison. All results in Table 2 are therefore reported on the matched $n=511$ subset.

F Timing Breakdown

Table 5 reports per-method timing statistics supplementing Table 2. Forward time dominates in all cases; constraint overhead is a small fraction for Dgrammar and DPGrammar. LAVE’s higher mean time is driven by timeout instances (capped at 120 s); its median is lower because the majority of instances either succeed quickly or exhaust their resample budget early. Effective constraint time for DPGrammar is near zero on instances where no DP repair is triggered; the non-trivial mean reflects the small fraction of instances that require one or more Viterbi passes.

Metric	LAVE	Dgrammar	DPGrammar
Avg fwd passes	99.2	104.7	99.4
Constraint overhead (%)	11.95	24.25	12.67
Eff. constraint (%)	—	20.43	1.22
Per-token total (ms)	111.5	100.7	71.9
Per-token constraint (ms)	6.64	8.26	0.30
Avg mask compute (ms)	—	427.2	577.0
Avg forward total (ms)	3951	4612	4229

Table 5: Detailed timing breakdown on jsb_medium ($n=511$ matched instances). Effective constraint % measures time spent in grammar-check calls that were on the critical path; it is not available for LAVE, which lacks the same instrumentation.

Algorithm 1 DPGrammar: Viterbi segment repair over a violated span (dp_fix_prefix)

Require: Violated span $[c, s]$; log-probabilities $\ell \in \mathbb{R}^{L \times |V|}$ from the current forward; grammar matcher $\mathcal{M}$ at the DFA state just before c ; beam width K

Ensure: Replacement pairs $\{(p_i, t_i)\}$ maximising $\sum \log p(t_i)$ subject to grammar, or FAIL

```

1:  $P \leftarrow [c, c+1, \dots, s-1]$ 
2:  $states \leftarrow \{\text{key}(\mathcal{M}) \mapsto (\mathcal{M}, 0)\}$  ▷ DFA-state  $\rightarrow$  (matcher, cumulative log-prob)
3:  $back \leftarrow \{\}$  ▷ backtrack table
4: for  $i \leftarrow 0$  to  $|P| - 1$  do
5:    $T_K \leftarrow \text{top-K}(\ell[P[i]], k=K)$ 
6:    $winners \leftarrow \{\}$ 
7:   for each (prev_key,  $(\mathcal{M}_p, \sigma_p)$ ) in  $states$  do
8:     for each  $(t, \ell_t)$  in  $T_K$  do
9:       if  $\mathcal{M}_p.\text{try\_consume}([t]) < 1$  then
10:         $\mathcal{M}_p.\text{rollback}(0)$ ; continue
11:       end if
12:        $\sigma' \leftarrow \sigma_p + \ell_t$ ;  $k' \leftarrow \text{key}(\mathcal{M}_p)$ 
13:        $\mathcal{M}_p.\text{rollback}(1)$ 
14:       if  $k' \notin winners$  or  $winners[k'].\sigma < \sigma'$  then
15:         $winners[k'] \leftarrow (\text{prev\_key}, t, \sigma')$  ▷ Viterbi merge
16:       end if
17:     end for
18:   end for
19:   if  $winners = \emptyset$  then return FAIL
20:   end if
21:    $next \leftarrow \{\}$ 
22:   for each  $(k', (\text{prev\_key}, t, \sigma'))$  in  $winners$  do ▷ Phase 2: clone once per surviving state
23:      $\mathcal{M}' \leftarrow states[\text{prev\_key}].\mathcal{M}.\text{deep\_copy}()$ ;  $\mathcal{M}'.\text{try\_consume}([t])$ 
24:      $next[k'] \leftarrow (\mathcal{M}', \sigma')$ ;  $back[(i, k')] \leftarrow (\text{prev\_key}, t)$ 
25:   end for
26:    $states \leftarrow next$ 
27: end for
28:  $k^* \leftarrow \arg \max_k states[k].\sigma$ ;  $fixes \leftarrow []$ ;  $cur \leftarrow k^*$ 
29: for  $i \leftarrow |P| - 1$  downto 0 do
30:    $(\text{prev}, t) \leftarrow back[(i, cur)]$ 
31:   if  $t \neq x[P[i]]$  then  $fixes.append(P[i], t)$ 
32:   end if
33:    $cur \leftarrow \text{prev}$ 
34: end for
35: return reverse( $fixes$ )

```
